

AI Ethics Critical Thinking Assignment

Case Study: Wrongful Arrest of Robert Williams via Facial Recognition (2020)

1. Case Study Selection and Summary

In January 2020, Robert Williams was wrongfully arrested by the Detroit Police Department after facial recognition software falsely identified his driver's license photo as a match for a shoplifting suspect captured on grainy security footage. Williams was arrested and handcuffed in front of his wife and two young daughters and detained for approximately 30 hours. Detectives did not perform adequate independent verification before obtaining an arrest warrant. The case highlighted the risks of using facial recognition technology without sufficient human verification.

2. Critical Analysis

Issue 1: Algorithmic Bias and Racial Disparity in Facial Recognition Accuracy

Why I chose this issue: I chose algorithmic bias because facial recognition algorithms have demonstrated higher false-positive rates for some demographic groups, particularly Black individuals. Williams's case provides a real-world example of how these technical disparities can create serious consequences when facial recognition is used in law enforcement.

My thought process: My first instinct was to treat this as a one-off software glitch, but looking closer, I realized that facial recognition algorithms are trained on datasets that underrepresent darker skin tones, which produces systematically higher error rates for Black individuals, making the problem more than an isolated technical failure. This reframed the issue for me from simply "the software glitched" to recognizing that a technology with known disparities in accuracy can predictably cause greater harm to certain populations, especially when used in law enforcement.

Potential solution: Regulators should mandate independent, demographically disaggregated testing before biometric systems are deployed in law enforcement. Police agencies should establish minimum accuracy and false-positive thresholds across relevant demographic groups rather than relying only on overall system accuracy.

Issue 2: Overreliance on AI Output as Sufficient Evidence

Why I chose this issue: Detectives never questioned Williams before arresting him, the algorithm's suggestion, paired with a single uncorroborated photo lineup identification, was treated as enough to justify an arrest. This highlights the gap between the limited evidence provided by the facial recognition result and the level of confidence investigators placed in its accuracy.

My thought process: I originally attributed Williams's wrongful arrest primarily to negligent police procedure, but examining the case highlighted the broader problem of automation bias, where people may place too much trust in AI-generated recommendations. This can lead investigators to overlook contradictory information or skip standard verification procedures, showing that the ethical concern is not only whether the algorithm is accurate but also whether humans understand its limitations and maintain appropriate evidentiary standards.

Potential solution: Mandate that facial recognition results be treated strictly as an unverified investigative lead, never as probable cause on their own, with binding department policy requiring independent corroborating evidence, such as an alibi check, before any arrest.

Issue 3: Accountability and Institutional Recourse

Why I chose this issue: It took Williams and the ACLU more than four years of litigation, from his 2020 arrest to a 2024 settlement, to secure accountability and policy change from the Detroit Police Department. This length of time highlights how little built-in recourse exists for someone harmed by an AI system.

My thought process: At first, I thought accountability just meant identifying who made the mistake and clearing Williams's record. But the more difficult question is what happens after the mistake: who bears the burden of pursuing correction, how long it takes, and whether the institution changes its practices or simply settles the individual case. That shifted my thinking toward accountability as an ongoing structural obligation rather than a one-time apology.

Potential solution: Build automatic reporting and independent review requirements into any law enforcement use of facial recognition, so that a false-positive match triggers a mandatory audit

and policy reassessment, rather than relying on the wrongly accused person to spend years pursuing a lawsuit.

3. Reflection on Critical Thinking

Working through this case shifted my thinking from viewing facial recognition failures as isolated software bugs to seeing them as predictable outcomes of how these systems are designed, tested, and deployed within existing institutions. I started out assuming the main fix was simply "improve the algorithm's accuracy." By the end, I recognized that even a highly accurate system can still cause serious harm if it's treated as infallible by the people using it, and if there's no fast, structural path to accountability when it fails. This changed my sense of where responsibility sits, not only with the engineers building the model, but with the institutions deciding how much weight to give its output. Going forward, this will push me to ask not just "how accurate is this system?" but "what happens to the people harmed when it's wrong, and how quickly can that be corrected?"

(Note: I used AI (ChatGPT & Claude) assistance to gather and verify factual details about this incident; all analysis and reflections are my own.)

References

American Civil Liberties Union. (n.d.). *Williams v. City of Detroit*.

<https://www.aclu.org/cases/williams-v-city-of-detroit-face-recognition-false-arrest>

NPR. (2020, June 24). '*The computer got it wrong*': How facial recognition led to false arrest of Black man. <https://www.npr.org/2020/06/24/882683463/the-computer-got-it-wrong-how-facial-recognition-led-to-a-false-arrest-in-michig>